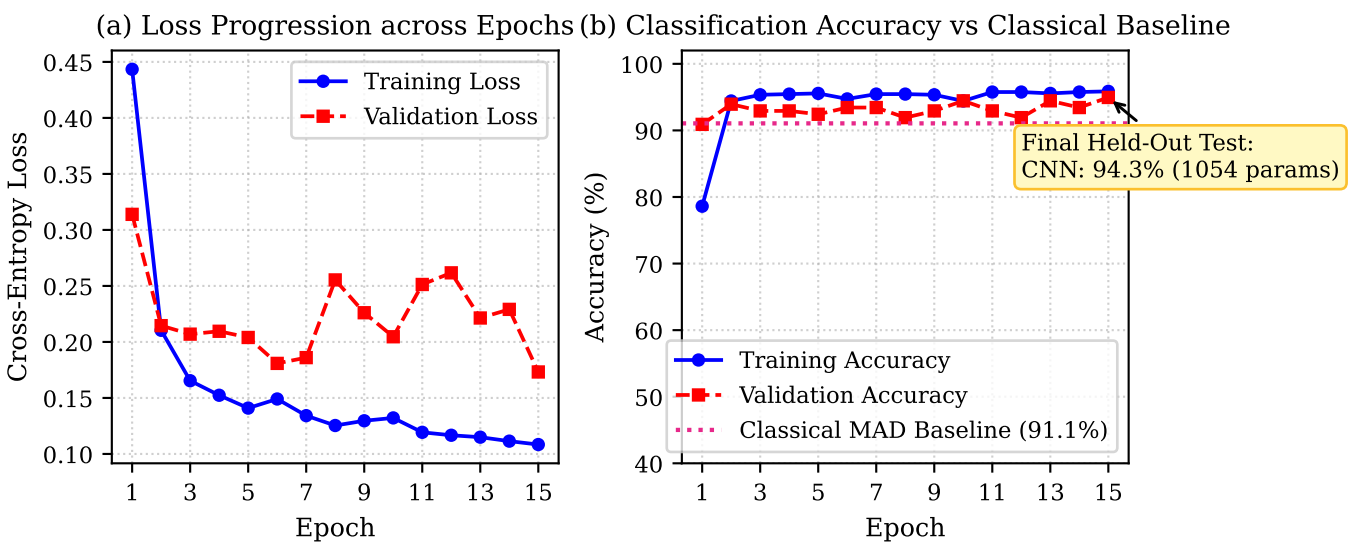
(a) Loss Progression across Epochs
(b) Classification Accuracy vs Classical Baseline
Cross-Entropy Loss
Epoch
Training Loss
Validation Loss
Accuracy (%)
Training Accuracy
Validation Accuracy
Classical MAD Baseline (91.1%)
Final Held-Out Test:
CNN: 94.3% (1054 params)